Question #580 *Topic 1*

A company uses locally attached storage to run a latency-sensitive application on premises. The company is using a lift and shift method to move the application to the AWS Cloud. The company does not want to change the application architecture.

Which solution will meet these requirements MOST cost-effectively?

A. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for Lustre file system to run the application.

B. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP2 volume to run the application.

C. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for OpenZFS file system to run the application.

D. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP3 volume to run the application.

Correct Answer: *B*

Community vote distribution

D (100%)

bojila 5 days, 14 hours ago
GP3 is the lastest version
upvoted 1 times

Hades2231 1 week, 1 day ago
Selected Answer: D
GP3 is the lastest version, and it is cost effective
upvoted 1 times

Guru4Cloud 2 weeks, 2 days ago
Selected Answer: D
GP3 is preferable over GP2, FSx for Lustre, and FSx for OpenZFS is clear and convincing:

GP3 offers identical latency performance to GP2 at a lower price point.
FSx options are higher performance but more expensive and require application changes.
GP3 aligns better with lift and shift needs as a directly attached block storage volume.
upvoted 1 times

taustin2 3 weeks, 3 days ago
Selected Answer: D
Migrate your Amazon EBS volumes from gp2 to gp3 and save up to 20% on costs.
upvoted 1 times

Vadbro7 3 weeks, 4 days ago
Y not gp2
upvoted 1 times

Ale1973 4 weeks ago
Selected Answer: D
My rational:
Options A y C are based on autoscaling-group and no make sense for me on this scenary.
Then, use Amazon EBS is the solution and GP2 or GP3 is the question.
Requirement requires the most COST effective solution, then, I choose GP3
upvoted 1 times

Question #581 *Topic 1*

A company runs a stateful production application on Amazon EC2 instances. The application requires at least two EC2 instances to always be running.

A solutions architect needs to design a highly available and fault-tolerant architecture for the application. The solutions architect creates an Auto Scaling group of EC2 instances.

Which set of additional steps should the solutions architect take to meet these requirements?

A. Set the Auto Scaling group's minimum capacity to two. Deploy one On-Demand Instance in one Availability Zone and one On-Demand Instance in a second Availability Zone.

B. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two On-Demand Instances in a second Availability Zone.

C. Set the Auto Scaling group's minimum capacity to two. Deploy four Spot Instances in one Availability Zone.

D. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two Spot Instances in a second Availability Zone.

Correct Answer: *D*

Community vote distribution

B (80%) A (20%)

luiscc Highly Voted 1 month, 1 week ago

Selected Answer: B

By setting the Auto Scaling group's minimum capacity to four, the architect ensures that there are always at least two running instances. Deploying two On-Demand Instances in each of two Availability Zones ensures that the application is highly available and fault-tolerant. If one Availability Zone becomes unavailable, the application can still run in the other Availability Zone.

upvoted 8 times

Guru4Cloud Most Recent 2 weeks, 2 days ago

Selected Answer: B

° Minimum of 4 ensures at least 2 instances are always running in each AZ, meeting the HA requirement.
° On-Demand instances provide consistent performance and availability, unlike Spot.
° Spreading across 2 AZs adds fault tolerance, protecting from AZ failure.

upvoted 2 times

Ale1973 4 weeks, 1 day ago

Selected Answer: A

My rational is: Highly available = 2 AZ, and then 2 EC2 instances always running is 1 EC2 in each AZ. If an entire AZ fails, SacalinGroup deploy the minimun instances (2) on the running AZ

upvoted 3 times

darkknight23 4 weeks, 1 day ago

Selected Answer: B

While Spot Instances can be used to reduce costs, they might not provide the same level of availability and guaranteed uptime that On-Demand Instances offer. So I will go with B and not D.

upvoted 1 times

Sat897 1 month ago

Selected Answer: B

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ.

upvoted 1 times

Sat897 1 month ago

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ..

upvoted 1 times

Question #582 *Topic 1*

An ecommerce company uses Amazon Route 53 as its DNS provider. The company hosts its website on premises and in the AWS Cloud. The company's on-premises data center is near the us-west-1 Region. The company uses the eu-central-1 Region to host the website. The company wants to minimize load time for the website as much as possible.

Which solution will meet these requirements?

A. Set up a geolocation routing policy. Send the traffic that is near us-west-1 to the on-premises data center. Send the traffic that is near eu-central-1 to eu-central-1.

B. Set up a simple routing policy that routes all traffic that is near eu-central-1 to eu-central-1 and routes all traffic that is near the on-premises datacenter to the on-premises data center.

C. Set up a latency routing policy. Associate the policy with us-west-1.

D. Set up a weighted routing policy. Split the traffic evenly between eu-central-1 and the on-premises data center.

Correct Answer: *A*

Community vote distribution

A (100%)

Hades2231 1 week, 1 day ago

Selected Answer: A

Geolocation is the key word

upvoted 1 times

lemur88 1 week, 4 days ago

Selected Answer: A

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html

upvoted 1 times

Guru4Cloud 2 weeks, 2 days ago

Selected Answer: A

The key reasons are:

Geolocation routing allows you to route users to the closest endpoint based on their geographic location. This will provide the lowest latency.
Routing us-west-1 traffic to the on-premises data center minimizes latency for those users since it is also located near there.
Routing eu-central-1 traffic to the eu-central-1 AWS region minimizes latency for users nearby.
This achieves routing users to the closest endpoint on a geographic basis to optimize for low latency.

upvoted 2 times

PLN6302 1 week, 5 days ago

why can't be the option C

upvoted 1 times

lemur88 1 week, 4 days ago

You cannot associate the policy to us-west-1 as the AWS account is in eu-central-1

upvoted 1 times

Question #583 *Topic 1*

A company has 5 PB of archived data on physical tapes. The company needs to preserve the data on the tapes for another 10 years for compliance purposes. The company wants to migrate to AWS in the next 6 months. The data center that stores the tapes has a 1 Gbps uplink internet connectivity.

Which solution will meet these requirements MOST cost-effectively?

A. Read the data from the tapes on premises. Stage the data in a local NFS storage. Use AWS DataSync to migrate the data to Amazon S3 Glacier Flexible Retrieval.

B. Use an on-premises backup application to read the data from the tapes and to write directly to Amazon S3 Glacier Deep Archive.

C. Order multiple AWS Snowball devices that have Tape Gateway. Copy the physical tapes to virtual tapes in Snowball. Ship the Snowball devices to AWS. Create a lifecycle policy to move the tapes to Amazon S3 Glacier Deep Archive.

D. Configure an on-premises Tape Gateway. Create virtual tapes in the AWS Cloud. Use backup software to copy the physical tape to the virtual tape.

Correct Answer: *C*

Community vote distribution

C (100%)

Hades2231 1 week, 1 day ago

Selected Answer: C

Ready for the exam tomorrow. Wish you guys all the best. BTW Snowball Device comes in handy when you need to move a huge amount of data but cant afford any bandwidth loss

upvoted 4 times

adeyinkaamole 1 week, 2 days ago

If you have made it to the end of the exam dump, you will definitely pass your exams in Jesus name. After over a year of Procrastination, I am finally ready to write my AWS Solutions Architect Exam. Thank you Exam Topics

upvoted 4 times

lemur88 1 week, 4 days ago

Selected Answer: C

Only thing that makes sense given the 1Gbps limitation

upvoted 1 times

Guru4Cloud 2 weeks, 2 days ago

Selected Answer: C

Option C is likely the most cost-effective solution given the large data size and limited internet bandwidth. The physical data transfer and integration with the existing tape infrastructure provides efficiency benefits that can optimize the cost.

upvoted 2 times

barracouto 3 weeks ago

Selected Answer: C

Went through this dump twice now. Exam is in about an hour. Will update with results.

upvoted 1 times

Vaishali12 2 weeks ago

how was ur exam?
was these dump que helpful?

upvoted 1 times

riccardoto 3 weeks, 6 days ago

Finished the dump today - taking my exam tomorrow :-) Wish me luck!

upvoted 2 times

Ale1973 4 weeks, 1 day ago

My rational: question is about which solution will meet these requirements MOST cost-effectively, not MOST time or effectively, then, my response is D (using Tape Gateways)

upvoted 2 times

D10SJoker 1 month ago

Selected Answer: C

For me it's C

upvoted 1 times

PrincePazol 1 month ago

Selected Answer: C

Taking my exams today

upvoted 1 times

mrsoa 1 month ago

Selected Answer: C

C is the right answer, because we need atleast 1 year to transfer the data over the internet

upvoted 2 times

Deepakin96 1 month ago

Selected Answer: C

C is my answer

upvoted 2 times

Question #584 *Topic 1*

A company is deploying an application that processes large quantities of data in parallel. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to prevent groups of nodes from sharing the same underlying hardware.

Which networking solution meets these requirements?

A. Run the EC2 instances in a spread placement group.

B. Group the EC2 instances in separate accounts.

C. Configure the EC2 instances with dedicated tenancy.

D. Configure the EC2 instances with shared tenancy.

Correct Answer: *A*

Community vote distribution

A (100%)

Eminenza22 6 days, 16 hours ago

Selected Answer: A

Option A is the correct answer. It suggests running the EC2 instances in a spread placement group. This solution is cost-effective and requires minimal development effort .

upvoted 1 times

Eminenza22 5 days, 18 hours ago

The placement group reduces the risk of simultaneous failures by spreading the instances across distinct underlying hardware

upvoted 1 times

czyboi 1 week ago

Selected Answer: A

A spread placement group is a group of instances that are each placed on distinct hardware.
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html

upvoted 1 times

Question #585 *Topic 1*

A solutions architect is designing a disaster recovery (DR) strategy to provide Amazon EC2 capacity in a failover AWS Region. Business requirements state that the DR strategy must meet capacity in the failover Region.

Which solution will meet these requirements?

A. Purchase On-Demand Instances in the failover Region.

B. Purchase an EC2 Savings Plan in the failover Region.

C. Purchase regional Reserved Instances in the failover Region.

D. Purchase a Capacity Reservation in the failover Region.

Correct Answer: *C*

Community vote distribution

D (75%) C (25%)

Eminenza22 2 days, 23 hours ago

Selected Answer: D

A regional Reserved Instance does not reserve capacity
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/reserved-instances-scope.html

upvoted 1 times

judyda 4 days, 2 hours ago

Selected Answer: D

reserved instances for price discount. need capacity reservation.

upvoted 2 times

gispankaj 5 days, 22 hours ago

Selected Answer: C

The Reserved Instance discount applies to instance usage within the instance family, regardless of size.

upvoted 1 times

ErnShm 5 days, 22 hours ago

D

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html

upvoted 1 times

Question #586 *Topic 1*

A company has five organizational units (OUs) as part of its organization in AWS Organizations. Each OU correlates to the five businesses that the company owns. The company's research and development (R&D) business is separating from the company and will need its own organization. A solutions architect creates a separate new management account for this purpose.

What should the solutions architect do next in the new management account?

A. Have the R&D AWS account be part of both organizations during the transition.

B. Invite the R&D AWS account to be part of the new organization after the R&D AWS account has left the prior organization.

C. Create a new R&D AWS account in the new organization. Migrate resources from the prior R&D AWS account to the new R&D AWS account.

D. Have the R&D AWS account join the new organization. Make the new management account a member of the prior organization.

Correct Answer: *C*

Community vote distribution

B (100%)

ErnShm 5 days, 21 hours ago

B
https://aws.amazon.com/blogs/mt/migrating-accounts-between-aws-organizations-with-consolidated-billing-to-all-features/

upvoted 1 times

gispankaj 5 days, 22 hours ago

Selected Answer: B

account can leave current organization and then join new organization.

upvoted 1 times

Question #587 *Topic 1*

A company is designing a solution to capture customer activity in different web applications to process analytics and make predictions. Customer activity in the web applications is unpredictable and can increase suddenly. The company requires a solution that integrates with other web applications. The solution must include an authorization step for security purposes.

Which solution will meet these requirements?

A. Configure a Gateway Load Balancer (GWLB) in front of an Amazon Elastic Container Service (Amazon ECS) container instance that stores the information that the company receives in an Amazon Elastic File System (Amazon EFS) file system. Authorization is resolved at the GWLB.

B. Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis data stream that stores the information that the company receives in an Amazon S3 bucket. Use an AWS Lambda function to resolve authorization.

C. Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis Data Firehose that stores the information that the company receives in an Amazon S3 bucket. Use an API Gateway Lambda authorizer to resolve authorization.

D. Configure a Gateway Load Balancer (GWLB) in front of an Amazon Elastic Container Service (Amazon ECS) container instance that stores the information that the company receives on an Amazon Elastic File System (Amazon EFS) file system. Use an AWS Lambda function to resolve authorization.

Correct Answer: *D*

Community vote distribution

C (100%)

Eminenza22 5 days, 17 hours ago

Selected Answer: C

https://docs.aws.amazon.com/lambda/latest/dg/services-kinesisfirehose.html

upvoted 1 times

ErnShm 5 days, 21 hours ago

C

authorizer is configured for the method. If it is, API Gateway calls the Lambda function. The Lambda function authenticates the caller by means such as the following: Calling out to an OAuth provider to get an OAuth access token

upvoted 1 times

gispankaj 5 days, 22 hours ago

Selected Answer: C

lambda authoriser seems to be logical solution.

upvoted 1 times

ralfj 6 days, 18 hours ago

Selected Answer: C

https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-use-lambda-authorizer.html

upvoted 4 times

Question #588 *Topic 1*

An ecommerce company wants a disaster recovery solution for its Amazon RDS DB instances that run Microsoft SQL Server Enterprise Edition. The company's current recovery point objective (RPO) and recovery time objective (RTO) are 24 hours.

Which solution will meet these requirements MOST cost-effectively?

A. Create a cross-Region read replica and promote the read replica to the primary instance.

B. Use AWS Database Migration Service (AWS DMS) to create RDS cross-Region replication.

C. Use cross-Region replication every 24 hours to copy native backups to an Amazon S3 bucket.

D. Copy automatic snapshots to another Region every 24 hours.

Correct Answer: *B*

Community vote distribution

D (100%)

Eminenza22 2 days, 17 hours ago

Selected Answer: D

This is the most cost-effective solution because it does not require any additional AWS services. Amazon RDS automatically creates snapshots of your DB instances every hour. You can copy these snapshots to another Region every 24 hours to meet your RPO and RTO requirements.

The other solutions are more expensive because they require additional AWS services. For example, AWS DMS is a more expensive service than AWS RDS.

upvoted 1 times

TiagueteVital 4 days, 11 hours ago

Selected Answer: D

Snapshots are always a cost-efficience way to have a DR plan.

upvoted 2 times

Question #589 *Topic 1*

A company runs a web application on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer that has sticky sessions enabled. The web server currently hosts the user session state. The company wants to ensure high availability and avoid user session state loss in the event of a web server outage.

Which solution will meet these requirements?

A. Use an Amazon ElastiCache for Memcached instance to store the session data. Update the application to use ElastiCache for Memcached to store the session state.

B. Use Amazon ElastiCache for Redis to store the session state. Update the application to use ElastiCache for Redis to store the session state.

C. Use an AWS Storage Gateway cached volume to store session data. Update the application to use AWS Storage Gateway cached volume to store the session state.

D. Use Amazon RDS to store the session state. Update the application to use Amazon RDS to store the session state.

Correct Answer: *D*

Community vote distribution

B (100%)

gispankaj 5 days, 22 hours ago

Selected Answer: B

redis is correct since it provides high availability and data persistance

upvoted 1 times

Eminenza22 6 days, 16 hours ago

Selected Answer: B

B is the correct answer. It suggests using Amazon ElastiCache for Redis to store the session state. Update the application to use ElastiCache for Redis to store the session state. This solution is cost-effective and requires minimal development effort.

upvoted 2 times

czyboi 1 week ago

Selected Answer: B

high availability => use redis instead of Elastich memcache

upvoted 2 times

Question #590 *Topic 1*

A company migrated a MySQL database from the company's on-premises data center to an Amazon RDS for MySQL DB instance. The company sized the RDS DB instance to meet the company's average daily workload. Once a month, the database performs slowly when the company runs queries for a report. The company wants to have the ability to run reports and maintain the performance of the daily workloads.

Which solution will meet these requirements?

A. Create a read replica of the database. Direct the queries to the read replica.

B. Create a backup of the database. Restore the backup to another DB instance. Direct the queries to the new database.

C. Export the data to Amazon S3. Use Amazon Athena to query the S3 bucket.

D. Resize the DB instance to accommodate the additional workload.

Correct Answer: *A*

Community vote distribution

A (100%)

Eminenza22 2 days, 17 hours ago

Selected Answer: A

This is the most cost-effective solution because it does not require any additional AWS services. A read replica is a copy of a database that is synchronized with the primary database. You can direct the queries for the report to the read replica, which will not affect the performance of the daily workloads

upvoted 1 times

TiagueteVital 4 days, 11 hours ago

Selected Answer: A

Clearly the right choice, with a read replica all the queries needed for a report are done in the replica, leaving the primary on best perfomance for write

upvoted 1 times

Question #591 *Topic 1*

A company runs a container application by using Amazon Elastic Kubernetes Service (Amazon EKS). The application includes microservices that manage customers and place orders. The company needs to route incoming requests to the appropriate microservices.

Which solution will meet this requirement MOST cost-effectively?

A. Use the AWS Load Balancer Controller to provision a Network Load Balancer.

B. Use the AWS Load Balancer Controller to provision an Application Load Balancer.

C. Use an AWS Lambda function to connect the requests to Amazon EKS.

D. Use Amazon API Gateway to connect the requests to Amazon EKS.

Correct Answer: *C*

Community vote distribution

D (100%)

Eminenza22 5 days, 17 hours ago

Selected Answer: D

https://aws.amazon.com/blogs/containers/microservices-development-using-aws-controllers-for-kubernetes-ack-and-amazon-eks-blueprints/

upvoted 1 times

ralfj 6 days, 18 hours ago

Selected Answer: D

https://aws.amazon.com/blogs/containers/integrate-amazon-api-gateway-with-amazon-eks/

upvoted 1 times

Question #592 *Topic 1*

A company uses AWS and sells access to copyrighted images. The company's global customer base needs to be able to access these images quickly. The company must deny access to users from specific countries. The company wants to minimize costs as much as possible.

Which solution will meet these requirements?

A. Use Amazon S3 to store the images. Turn on multi-factor authentication (MFA) and public bucket access. Provide customers with a link to the S3 bucket.

B. Use Amazon S3 to store the images. Create an IAM user for each customer. Add the users to a group that has permission to access the S3 bucket.

C. Use Amazon EC2 instances that are behind Application Load Balancers (ALBs) to store the images. Deploy the instances only in the countries the company services. Provide customers with links to the ALBs for their specific country's instances.

D. Use Amazon S3 to store the images. Use Amazon CloudFront to distribute the images with geographic restrictions. Provide a signed URL for each customer to access the data in CloudFront.

Correct Answer: *C*

Community vote distribution

D (100%)

hubbabubba 4 days, 16 hours ago

Selected Answer: D

answer is D

upvoted 1 times

Eminenza22 5 days, 17 hours ago

Selected Answer: D

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html

upvoted 1 times

ralfj 6 days, 18 hours ago

Selected Answer: D

Use Cloudfront and geographic restriction

upvoted 3 times

Question #593 *Topic 1*

A solutions architect is designing a highly available Amazon ElastiCache for Redis based solution. The solutions architect needs to ensure that failures do not result in performance degradation or loss of data locally and within an AWS Region. The solution needs to provide high availability at the node level and at the Region level.

Which solution will meet these requirements?

A. Use Multi-AZ Redis replication groups with shards that contain multiple nodes.

B. Use Redis shards that contain multiple nodes with Redis append only files (AOF) turned on.

C. Use a Multi-AZ Redis cluster with more than one read replica in the replication group.

D. Use Redis shards that contain multiple nodes with Auto Scaling turned on.

Correct Answer: *A*

Community vote distribution

B (67%) A (33%)

hubbabubba 4 days, 16 hours ago

Selected Answer: A

Hate to say this, but I read the two docs linked below, and I still think the answer is A. Turning on AOF helps in data persistence after failure, but it does nothing for availability unless you use Multi-AZ replica groups.

upvoted 1 times

Eminenza22 5 days, 17 hours ago

Selected Answer: B

https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/FaultTolerance.html

upvoted 1 times

ralfj 6 days, 18 hours ago

Selected Answer: B

https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/RedisAOF.html

upvoted 1 times

Question #594 *Topic 1*

A company plans to migrate to AWS and use Amazon EC2 On-Demand Instances for its application. During the migration testing phase, a technical team observes that the application takes a long time to launch and load memory to become fully productive.

Which solution will reduce the launch time of the application during the next testing phase?

A. Launch two or more EC2 On-Demand Instances. Turn on auto scaling features and make the EC2 On-Demand Instances available during the next testing phase.

B. Launch EC2 Spot Instances to support the application and to scale the application so it is available during the next testing phase.

C. Launch the EC2 On-Demand Instances with hibernation turned on. Configure EC2 Auto Scaling warm pools during the next testing phase.

D. Launch EC2 On-Demand Instances with Capacity Reservations. Start additional EC2 instances during the next testing phase.

Correct Answer: *C*

Community vote distribution

C (100%)

ralfj 6 days, 19 hours ago

Selected Answer: C

just use hibernation option so you won't load the full EC2 Instance

upvoted 1 times

Question #595 *Topic 1*

A company's applications run on Amazon EC2 instances in Auto Scaling groups. The company notices that its applications experience sudden traffic increases on random days of the week. The company wants to maintain application performance during sudden traffic increases.

Which solution will meet these requirements MOST cost-effectively?

A. Use manual scaling to change the size of the Auto Scaling group.

B. Use predictive scaling to change the size of the Auto Scaling group.

C. Use dynamic scaling to change the size of the Auto Scaling group.

D. Use schedule scaling to change the size of the Auto Scaling group.

Correct Answer: *C*

Community vote distribution

C (100%)

ralfj 6 days, 19 hours ago

Selected Answer: C

Dynamic Scaling – This is yet another type of Auto Scaling in which the number of EC2 instances is changed automatically depending on the signals received. Dynamic Scaling is a good choice when there is a high volume of unpredictable traffic.

https://www.developer.com/web-services/aws-auto-scaling-types-best-practices/#:~:text=Dynamic%20Scaling%20%E2%80%93%20This%20is%20yet,high%20volume%20of%20unpredictable%20traffic.

upvoted 1 times

Question #596 *Topic 1*

An ecommerce application uses a PostgreSQL database that runs on an Amazon EC2 instance. During a monthly sales event, database usage increases and causes database connection issues for the application. The traffic is unpredictable for subsequent monthly sales events, which impacts the sales forecast. The company needs to maintain performance when there is an unpredictable increase in traffic.

Which solution resolves this issue in the MOST cost-effective way?

A. Migrate the PostgreSQL database to Amazon Aurora Serverless v2.

B. Enable auto scaling for the PostgreSQL database on the EC2 instance to accommodate increased usage.

C. Migrate the PostgreSQL database to Amazon RDS for PostgreSQL with a larger instance type.

D. Migrate the PostgreSQL database to Amazon Redshift to accommodate increased usage.

Correct Answer: *C*

Community vote distribution

A (100%)

TiagueteVital 4 days, 10 hours ago

Selected Answer: A

A to autoscaling

upvoted 2 times

manOfThePeople 6 days, 13 hours ago

Answer is A.
Aurora Serverless v2 got autoscaling, highly available and cheaper when compared to the other options.

upvoted 1 times

anikety123 6 days, 15 hours ago

Selected Answer: A

The correct answer is A

upvoted 1 times

Question #597 *Topic 1*

A company hosts an internal serverless application on AWS by using Amazon API Gateway and AWS Lambda. The company's employees report issues with high latency when they begin using the application each day. The company wants to reduce latency.

Which solution will meet these requirements?

A. Increase the API Gateway throttling limit.

B. Set up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day.

C. Create an Amazon CloudWatch alarm to initiate a Lambda function as a target for the alarm at the beginning of each day.

D. Increase the Lambda function memory.

Correct Answer: *B*

Community vote distribution

B (100%)

Eminenza22 6 days, 16 hours ago

Selected Answer: B

B option setting up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day. This solution is cost-effective and requires minimal development effort.

upvoted 1 times

oayoade 1 week ago

Selected Answer: B

https://aws.amazon.com/blogs/compute/scheduling-aws-lambda-provisioned-concurrency-for-recurring-peak-usage/

upvoted 2 times

Question #598 *Topic 1*

A research company uses on-premises devices to generate data for analysis. The company wants to use the AWS Cloud to analyze the data. The devices generate .csv files and support writing the data to an SMB file share. Company analysts must be able to use SQL commands to query the data. The analysts will run queries periodically throughout the day.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose three.)

A. Deploy an AWS Storage Gateway on premises in Amazon S3 File Gateway mode.

B. Deploy an AWS Storage Gateway on premises in Amazon FSx File Gateway made.

C. Set up an AWS Glue crawler to create a table based on the data that is in Amazon S3.

D. Set up an Amazon EMR cluster with EMR File System (EMRFS) to query the data that is in Amazon S3. Provide access to analysts.

E. Set up an Amazon Redshift cluster to query the data that is in Amazon S3. Provide access to analysts.

F. Setup Amazon Athena to query the data that is in Amazon S3. Provide access to analysts.

Correct Answer: *CEF*

Community vote distribution

ACF (100%)

Eminenza22 5 days, 18 hours ago

Selected Answer: ACF

https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-format-csv-home.html
https://aws.amazon.com/blogs/aws/amazon-athena-interactive-sql-queries-for-data-in-amazon-s3/
https://aws.amazon.com/storagegateway/faqs/

upvoted 1 times

anikety123 6 days, 16 hours ago

Selected Answer: ACF

It should be ACF

upvoted 1 times

ralfj 6 days, 19 hours ago

Selected Answer: ACF

ACF use S3 File Gateway, Use Glue and Use Athena

upvoted 1 times

Question #599 *Topic 1*

A company wants to use Amazon Elastic Container Service (Amazon ECS) clusters and Amazon RDS DB instances to build and run a payment processing application. The company will run the application in its on-premises data center for compliance purposes.

A solutions architect wants to use AWS Outposts as part of the solution. The solutions architect is working with the company's operational team to build the application.

Which activities are the responsibility of the company's operational team? (Choose three.)

A. Providing resilient power and network connectivity to the Outposts racks

B. Managing the virtualization hypervisor, storage systems, and the AWS services that run on Outposts

C. Physical security and access controls of the data center environment

D. Availability of the Outposts infrastructure including the power supplies, servers, and networking equipment within the Outposts racks

E. Physical maintenance of Outposts components

F. Providing extra capacity for Amazon ECS clusters to mitigate server failures and maintenance events

Correct Answer: *ACE*

Community vote distribution

ACE (57%) ACF (29%) 14%

ibu007 3 days, 20 hours ago

Selected Answer: ACE

My exam is tomorrow. thank you all for the answers and links.

upvoted 3 times

Eminenza22 5 days, 18 hours ago

Selected Answer: ACF

A - With Outposts, you are responsible for providing resilient power and network connectivity to the Outpost racks to meet your availability requirements for workloads running on Outposts.

upvoted 2 times

Eminenza22 5 days, 18 hours ago

C - With AWS Outposts, you are responsible for the physical security and access controls of the data center environment.

upvoted 2 times

Eminenza22 5 days, 18 hours ago

F - Since Outpost capacity is finite and determined by the size and number of racks AWS installs at your site, "you" must decide how much EC2, EBS, and S3 on Outposts capacity "you" need to run your initial workloads, accommodate future growth, and to provide extra capacity to mitigate server failures and maintenance events

https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html

upvoted 2 times

ralfj 6 days, 19 hours ago

Selected Answer: ACD

https://docs.aws.amazon.com/outposts/latest/userguide/outposts-requirements.html

upvoted 1 times

ralfj 6 days, 11 hours ago

missed clicked, Should be ACE

upvoted 1 times

SOMEONE1675 1 week ago

Selected Answer: ACE

best answer

upvoted 1 times